

1. Policy corpus and ground truth

seeded generator: 7 risk categories $\times$ 5 tiers

7,000 records, labelled by gating matrix and by rubric

rubric reads the policy?

no

yes

not independent

2. Specification under test

monotone tier policy $a(c, t)$, nine capabilities

yes

monotone?

no

reject policy edit

enumerate 8,640 region-complete responses $\times$ 5 tiers

invariants hold?

yes

no

specification defect

3. Real children's language

three essay corpora, 31,307 essays

five features per 50-word window; duplicates removed

learned estimators, trained on two corpora

prompt held out?

yes

no

topic confound: report separately

4. Measurement

frameworks, ablations, tuned baselines

Wilson and bootstrap intervals, 20 seeds

at the ceiling?

yes

no

report recoverable headroom

verdict, with error direction